

Assessment Design for AI-Enabled Classrooms

Authentic assessment and academic integrity when every learner has generative AI

Generative AI is now part of the environment every learner works in. The question is no longer whether students will use it, but whether an assessment still measures something real when they do. The goal is authentic skill development. This document sets out how I design assessment that stays valid in an AI-enabled classroom, by design rather than by surveillance.

POSTURE

Three honest positions on AI use

Every assessment should declare which one it takes, so expectations are explicit and integrity is a shared agreement rather than a guessing game.

POSTURE	WHAT IT MEANS	FITS WHEN
AI-free	No AI. The point is to demonstrate the underlying capability unaided.	Foundational skills a professional must hold without a tool.
AI-assisted	AI is allowed and disclosed. The learner is assessed on judgment, not keystrokes.	Most applied work. Mirrors the real workplace.
AI-required	The task cannot be done well without AI, and using it well is the skill.	Prompting, evaluation, and oversight of AI output.

Integrity is not won by catching cheaters. It is designed in by asking for evidence a model cannot produce on the learner's behalf: their reasoning, their process, their application to a context only they hold.

DESIGN MOVES

What keeps an assessment authentic

- **Assess the process, not only the product.** Ask for the draft, the dead ends, the decision log. A finished artifact is easy to generate; a defensible path to it is not.
- **Anchor in a context only the learner has.** Their workplace, their dataset, their patient population, their community. Specificity defeats the generic answer.
- **Require a live or spoken defense.** A short oral check confirms the thinking belongs to the person who submitted it.
- **Make AI use visible.** Ask learners to show what they prompted, what the model returned, and where their judgment overrode it. Disclosure becomes part of the work.
- **Assess critical evaluation of AI output.** Hand them a flawed AI answer and grade the critique. Spotting the error is the higher-order skill.

From AI-defeated to AI-resilient

Applied Data Literacy, Module 4. The original assessment asked learners to write a paragraph explaining why a data model was sound. A model answers that in seconds, so it measured nothing.

BEFORE	AFTER
Write a paragraph explaining why this schema is well designed.	Audit this deliberately flawed schema against the four business rules. Identify three failures, propose a fix for each, and record one place where you used AI and one place you chose not to, with your reasoning.
Output a model can generate. No evidence of the learner's thinking. No defensible record.	Application to a specific artifact, a critique task, a process record, and disclosed AI use. The skill is the judgment, and the judgment is visible.

INTEGRITY STANCE

The agreement, stated plainly

Authentic learning in an AI age rests on transparency, not policing. Learners are told which posture an assessment takes, are expected to disclose AI use honestly, and are graded on the thinking that is theirs. Critical thinking, not avoidance of the tool, is the outcome. An assessment designed this way is fairer, more like the work students are headed into, and far more durable than any detector.

Abstracted sample. Client names and proprietary content withheld. Structure and method are representative of production work.